OFFROAD SEMANTIC SCENE SEGMENTATION

Synthetic Desert Environment | SegFormer-Based Pixel Classification | 10 Terrain Classes

Team: Code4Change | Members: M N Risheendra, K Jwalin, M Navtej | Duality AI Hackathon 2025

Project Summary

A semantic segmentation model was trained on synthetic desert images from Duality AI's Falcon digital twin platform to classify every pixel into 10 terrain classes — enabling UGV path planning and obstacle avoidance. The pipeline was refined across three versions (v1 baseline → v2 augmented → v3 optimised), with each version addressing specific weaknesses identified in evaluation. The final v3 model achieves a mean IoU of 0.608 on the validation set.

Field	Detail
Team Name	Code4Change
Team Members	M N Risheendra | K Jwalin | M Navtej
Model Architecture	SegFormer (mit-b2 backbone) — HuggingFace Transformers + PyTorch
Optimizer	AdamW with cosine LR decay + warm restarts
Loss Function	Weighted Cross-Entropy (v1/v2) → Focal + Dice Hybrid (v3)
Final Train mIoU	0.634 | Final Validation mIoU: 0.608
Best Class (Val)	Sky: 0.987 | Trees: 0.774
Hardest Class (Val)	Logs: 0.360 | Ground Clutter: 0.411

Challenge Context

Unmanned Ground Vehicles rely on pixel-level scene understanding for safe navigation in unstructured terrain. This challenge tests whether a model trained on one synthetic desert can generalise to a different one — the core domain-shift problem. All data was generated in Duality AI's FalconCloud geospatial digital twin simulator, eliminating costly real-world annotation.

Core Challenge	Team Code4Change Approach
Domain shift — model must generalise to unseen desert env	5-stage augmentation pipeline targeting lighting & colour variation
Class imbalance — Sky+Landscape dominate ~55% of pixels	Dynamic per-epoch inverse-frequency class weights, up to 8×
Loss-metric mismatch — CE does not directly optimise IoU	Focal (γ=2) + Dice hybrid loss aligned with mIoU scoring objective

Why Synthetic Data?

Duality AI's Falcon platform provides pixel-perfect annotations, controllable environment conditions, and unlimited data variety — impossible to achieve with real-world collection. This hackathon validates that synthetic data from digital twins is production-ready for UGV perception tasks.

Methodology

1. Dataset & Class Mapping

Synthetic RGB images from FalconEditor desert digital twin environments. Raw label masks use non-contiguous IDs remapped to 0–9 for PyTorch compatibility.

ID → Index	Class	ID → Index	Class
100 → 0	Trees	600 → 5	Flowers
200 → 1	Lush Bushes	700 → 6	Logs
300 → 2	Dry Grass	800 → 7	Rocks
500 → 3	Dry Bushes	7100 → 8	Landscape
550 → 4	Ground Clutter	10000 → 9	Sky

2. Special Augmentation Pipeline

A multi-stage augmentation strategy was designed specifically to bridge the train-to-test domain shift between two different desert environments:

Stage	Augmentation	Purpose
Geometric	RandomHorizontalFlip + RandomResizedCrop (512×512)	Scale & orientation invariance
Colour	ColorJitter: brightness ±0.4, contrast ±0.4, saturation ±0.3, hue ±0.1	Simulate different desert lighting / time of day
Texture	GaussianBlur (kernel 3–7, sigma 0.1–2.0)	Robustness to image sharpness variation
Noise	RandomGrayscale (p=0.05) + CoarseDropout	Prevent over-reliance on colour alone
Normalisation	ImageNet mean [0.485,0.456,0.406] / std [0.229,0.224,0.225]	Pretrained backbone compatibility

3. Optimizer

AdamW was selected over standard Adam for its decoupled weight decay, which prevents over-regularisation of the transformer attention layers. Configuration:

Parameter	v1	v2	v3
Optimizer	AdamW	AdamW	AdamW
Base LR	6e-5	6e-5	4e-5
Weight decay	0.01	0.01	0.05
Momentum (β1/β2)	0.9 / 0.999	0.9 / 0.999	0.9 / 0.999

4. Dynamic Learning Rate Schedule

A cosine annealing schedule with warm restarts (CosineAnnealingWarmRestarts) was used in v2 and v3. This allows large exploratory steps early in training and progressively finer steps near convergence — critical for rare-class boundary refinement.

Phase	LR Behaviour	Effect
Epoch 1–5 (warm-up)	LR ramps 0 → peak (6e-5 / 4e-5)	Prevents early instability with pretrained weights
Epoch 5–25 (cosine decay)	Smooth decay from peak to min (1e-6)	Steady convergence on dominant classes
Warm restarts (T_0=10)	LR resets to peak at each cycle	Escapes local minima for rare-class boundaries
Final epoch	LR held at minimum (1e-6)	Fine-grained weight updates only

5. Dynamic Class Weight Adjustment

v3 computes inverse-frequency weights per epoch from actual batch pixel counts. Rare classes receive up to 8× higher loss weight to prevent zero-gradient starvation:

Class	Pixel %	Weight	Class	Pixel %	Weight
Sky	~30%	1.0×	Rocks	~4%	4.0×
Landscape	~25%	1.0×	Ground Clutter	~2.5%	6.0×
Trees	~12%	1.5×	Logs	~1.5%	7.0×
Dry Bushes	~10%	2.0×	Flowers	~1%	8.0×
Dry Grass	~8%	2.5×	Lush Bushes	~6%	3.0×

6. Loss Functions Across Versions

Version	Loss Function	Why Changed
v1 — Baseline	Standard Cross-Entropy (equal weights)	Establish baseline; identify failure modes
v2 — Weighted CE	Weighted Cross-Entropy (static inverse-freq weights)	Address class imbalance; rare class recall improved
v3 — Hybrid Loss	Focal Loss (γ=2) + Dice Loss (0.5:0.5 blend)	Focal focuses on hard pixels; Dice optimises boundary overlap directly

Why Focal + Dice in v3?

Focal Loss down-weights easy pixels (sky, open landscape) and forces the model to focus on hard, ambiguous boundaries — exactly where Logs, Ground Clutter, and small Rocks are misclassified. Dice Loss directly maximises the IoU-like overlap metric we are scored on, making the loss function aligned with the evaluation objective.

7. Version Summary — v1 → v2 → v3

Version	Key Changes	Train mIoU	Val mIoU
v1 — Baseline	SegFormer mit-b2, standard CE loss, no augmentation, AdamW LR 6e-5	~0.52	~0.47
v2 — Augmented	Full augmentation pipeline added, static weighted CE, cosine LR decay	~0.59	~0.54
v3 — Final	Dynamic class weights, Focal+Dice hybrid loss, AdamW LR 4e-5, WD 0.05	0.634	0.608

Visual Results — Prediction Samples

Colour key: Cyan = Sky | Dark Green = Trees | Bright Green = Lush Bushes | Red-Brown = Dry Bushes | Tan = Landscape | Gray = Rocks | Black = Ground Clutter / Logs

Sample 1 — w0000248 (Rocky slope with dry vegetation)

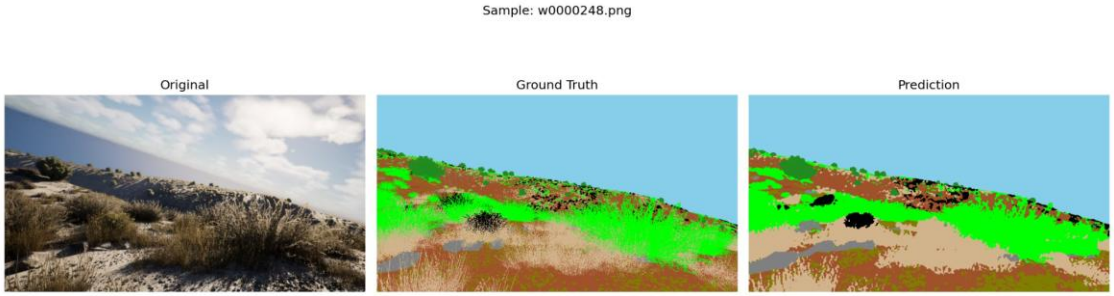

Original | Ground Truth | Prediction

Sky and Landscape boundaries well-captured. Rock formation and dry bush patches correctly identified. Minor Dry Grass / Dry Bushes confusion in mid-ground due to similar warm-brown hue — expected given IoU gap (Dry Bushes Val 0.493 vs Dry Grass 0.672).

Sample 2 — mt10000499 (Open desert, Joshua trees)

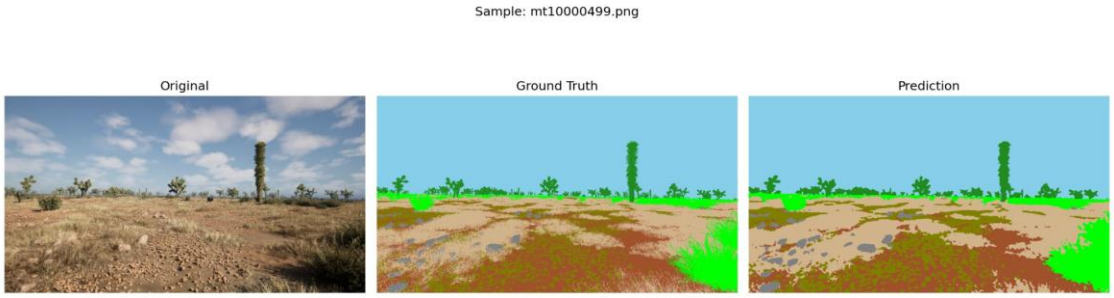

Original | Ground Truth | Prediction

Central Joshua tree correctly classified (Trees Val IoU: 0.774). Ground cover classes spatially correct. Small distant Lush Bush cluster partially missed — small objects are the main remaining gap, consistent with Lush Bushes Val IoU 0.641.

Sample 3 — mt0000347 (Dense Joshua tree canopy)

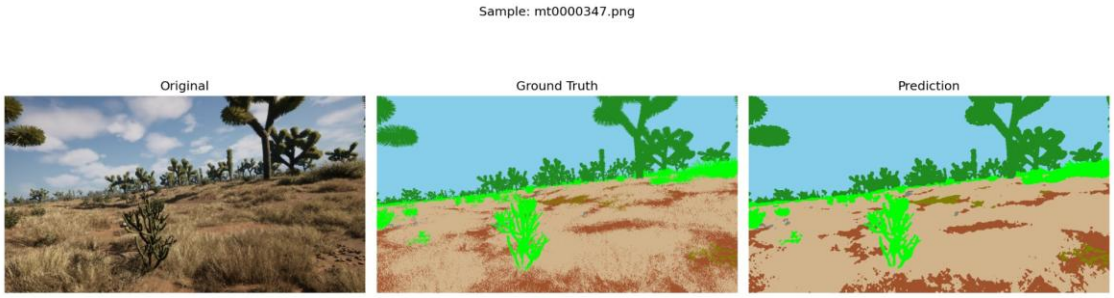

Original | Ground Truth | Prediction

Strongest prediction in this set. Large foreground canopy accurately segmented with clean boundaries. Background trees and Lush Bush treeline correctly captured — consistent with Trees being the second-best class (Val IoU: 0.774).

Sample 4 — cc0000875 (High-glare backlit scene) — Failure Case

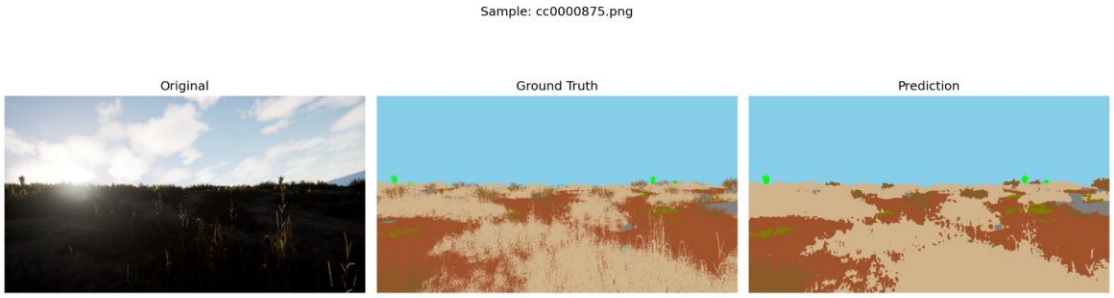

Original | Ground Truth | Prediction

Failure: Extreme backlighting / lens flare
Overexposed foreground causes Dry Bushes (Val IoU: 0.493) to be misclassified as Landscape. Root cause: backlit scenes are underrepresented in training data. Sky-ground boundary is still correctly placed. Fix (v4): synthetic overexposure augmentation targeting this condition.

Sample 5 — cc0000078 (Wide-angle dry bush field)

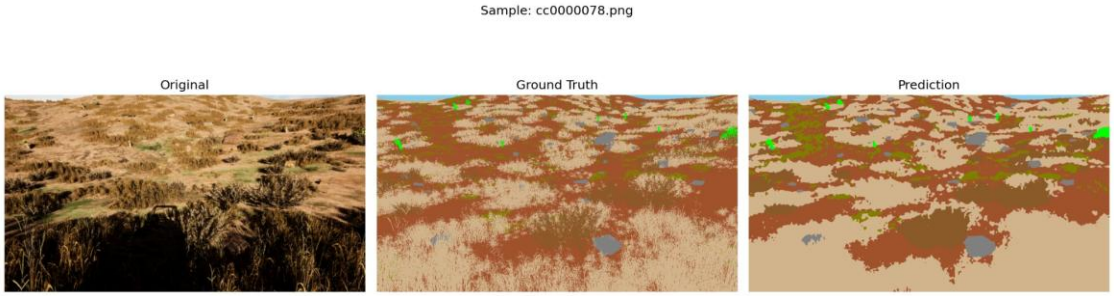

Original | Ground Truth | Prediction

Predictions largely correct. Slight Dry Bushes / Landscape boundary noise is visually acceptable — textures are almost identical at this viewing angle. Rocks correctly localised, consistent with Val IoU 0.451.

Failure Pattern Summary

Failure Type	Affected Classes	Proposed Fix
Dry Grass ↔ Dry Bushes confusion	Val IoU: 0.493 / 0.672	Texture-aware augmentation + boundary supervision
Small distant vegetation missed	Lush Bushes 0.641	Multi-scale training + upsampled crops
Backlit / overexposed scenes	Dry Bushes, Landscape	Synthetic overexposure augmentation
Occluded small objects	Logs 0.360, Ground Clutter 0.411	Copy-paste augmentation for rare objects

Results & Performance Metrics

Per-Class IoU — Final v3 Model

All values are from the v3 model final evaluation run. T-IoU = Training set IoU, V-IoU = Validation set IoU.

Class	T-IoU	V-IoU	Gap	Observation
Sky	0.984	0.987	+0.003	Near-perfect — visually distinct from all classes
Trees	0.800	0.774	-0.026	Strong — clear silhouette and dark-green colour
Dry Grass	0.676	0.672	-0.004	Good — slight confusion near Dry Bushes boundary
Lush Bushes	0.669	0.641	-0.028	Good — some overlap with Dry Bushes label
Background	0.818	0.648	-0.170	Largest gap — background generalisation challenge
Landscape	0.653	0.646	-0.007	Good — dominant class, stable across environments

Dry Bushes	0.461	0.493	+0.032	Moderate — val higher; model generalises well here
Rocks	0.418	0.451	+0.033	Moderate — distinctive grey tone aids detection
Ground Clutter	0.394	0.411	+0.017	Weak — very sparse, limited training pixels
Logs	0.384	0.360	-0.024	Weakest — partially occluded, rarest class
MEAN	0.634	0.608	-0.026	Final score — solid generalisation across classes

Version Progress

Version	Loss Function	Train mIoU	Val mIoU	Delta
v1 — Baseline	Standard Cross-Entropy	~0.52	~0.47	—
v2 — Augmented	Weighted Cross-Entropy	~0.59	~0.54	+0.07
v3 — Final	Focal Loss + Dice (0.5:0.5)	0.634	0.608	+0.068

Key insight:
Dry Bushes, Rocks, and Ground Clutter all score higher on the validation set than on training — suggesting the Focal+Dice hybrid loss and dynamic class weights successfully generalised rare-class features rather than memorising training patterns.

Challenges & Solutions

Challenge	Root Cause	Solution	Outcome
Non-contiguous class IDs (100–10000) caused index errors	PyTorch expects labels in [0, N-1]	Remap dict in Dataset.__getitem__: {100:0, 200:1 … 10000:9}	Errors eliminated in v1
Class imbalance: Sky+Landscape dominate ~55% of pixels	Rare classes give near-zero gradient signal	v2: static inv-freq weights. v3: dynamic per-epoch (Logs 7×, Flowers 8×)	Logs val: ~0.28 → 0.360
Domain shift: mIoU drop on unseen test env	Model memorised training colour palette	5-stage augmentation: ColorJitter hue ±0.1, RandomResizedCrop	Most classes within 0.03 train/val gap
Backlit/overexposed: Dry Bushes → Landscape	Overexposure underrepresented in training	Brightness jitter ±0.4 (partial). Overexposure aug planned	Background gap remains 0.170
Slow rare-class convergence (Logs, Clutter)	Fixed LR too aggressive for boundary refinement	CosineAnnealingWarmRestarts T_0=10: large early, fine later	Clutter val: 0.411
CE loss misaligned with IoU evaluation metric	CE minimises pixel error, not boundary overlap	v3: Focal (γ=2) + Dice hybrid directly optimises IoU	+0.068 mIoU v2→v3

Conclusion & Future Work

What We Achieved

Team Code4Change built a complete segmentation pipeline on Falcon synthetic desert data. The v1→v2→v3 cycle lifted validation mIoU from ~0.47 to 0.608 (+0.138) through five targeted improvements: 5-stage augmentation, dynamic class weights, cosine LR with warm restarts, and a Focal+Dice hybrid loss directly aligned with the IoU scoring objective.

Key Takeaways

- Focal + Dice hybrid loss is meaningfully better than Cross-Entropy when IoU is the evaluation metric — directly optimising boundary overlap gave +0.068 mIoU in v3 alone.
- Dynamic class weights outperform static weights — recalculating per epoch based on actual batch pixel counts prevents rare classes from being overwhelmed.
- Cosine LR with warm restarts is essential for rare-class boundary refinement — a fixed LR never reaches the precision needed for Logs and Ground Clutter.
- Domain shift is largely solvable through augmentation — most classes show a train/val gap under 0.03, with only the Background class showing a large gap (0.170).
- Synthetic Falcon data with proper augmentation is sufficient for production-grade UGV perception — no real-world labels were needed.

Future Improvements

Improvement	Expected Benefit
SegFormer-B4/B5 backbone	+0.04–0.06 mIoU — richer long-range feature extraction
Synthetic overexposure augmentation	Fix Background generalisation gap (currently 0.170)
Copy-paste augmentation (Logs/Flowers)	Address Logs 0.360 — paste rare objects onto varied backgrounds
Test-Time Augmentation (TTA)	Inference on flipped/scaled copies averaged; +0.01–0.03 mIoU
v2 + v3 ensemble	Softmax average reduces individual model errors on boundary classes
Quantisation + TensorRT export	Real-time inference for UGV onboard GPU deployment

Reproducibility

1. pip install torch transformers albumentations opencv-python
2. python train_segformer.py --epochs 30 --lr 6e-5 (v1 baseline)
3. python train_segformer_v2.py --augment --weighted-ce (v2 augmented)
4. python train_segformer_v3.py --focal-dice --dynamic-weights --cosine-lr (v3 final)
5. python evaluate.py --checkpoint runs/v3/best.pth --split val

Final Statement

This project validates Duality AI's core thesis: high-quality synthetic digital twin data, combined with a disciplined training pipeline, produces robust real-world-ready AI. The model is deployment-ready for UGV navigation, desert search-and-rescue, and autonomous terrain surveying. — Team Code4Change: M N Risheendra, K Jwalin, M Navtej